
Completely Meaningful Binary Encoding

November 5, 2023

Kees-Jan Hermans (kees.jan.hermans@gmail.com)

This treatise tries to describe a (useful) binary encoding of complex data types which cannot break the parser because of encoding related inconsistencies.

Contents

Contents	2
1 Premise	3
2 Problem Statement	3
3 Design Principles	4
4 Specification	4
5 Thoughts on Specific Forms of Sub-Encoding	5
5.1 Bigger type space encoding	5
5.2 String Encoding	6
5.3 Integer Encoding	6
5.4 Floating Point Number Encoding	6
6 Examples	6
7 Assessment	7
7.1 Advantages	7
7.2 Disadvantages	8

1 Premise

The premise of this treatise is that it must be possible to create a binary encoding format specification such that any random input deterministically produces a valid output. The use of this specification must allow for complex data structures in the output (for example, as offered by ASN.1, XML or JSON), and even useful ones (in other words, this is not just a joke).

2 Problem Statement

Binary encodings are often chosen because they are (supposedly) dense. However, when applied, they often contain at least some 'air' (superfluous or useless information). The problem with binary encodings is that this 'air' can also contain encoding mistakes. For example:

- When using TLV, they may contain types that aren't defined.
- When using TLV, they may contain lengths that are illegal. Either because they are zero, which the type does not support, or because they exceed the length of the message size, or indeed the size of the element that contains them.
- For example, in DER encoding it is possible to define lengths that could not possibly exist in our universe.
- For example, in Unicode, there exist illegal encodings, when the amount of bytes part of the definition of a character, aren't available, or when the resultant character value isn't specified.

Textual formats can convey complex data structures, but are also (very) error prone. For example:

- String escaping may lead to illegal escape sequences or breaking of string enclosing logic.
- Closing strings and brackets may be omitted or over-supplied.
- Closing tags - same thing.
- Illegal bytes in 7-bits clean formats.

etc.

All of these formats are problematic because we have to deal with errors during parsing. Errors in parsing are hardly ever meaningful, since we have no good way to recover from them (but they do introduce a lot of extra code). Error correcting codes can be provided with any input, this one being no exception.

Tail end recovery of both 'JSON'- or 'DER'-like encodings should be relatively easy (just wrap up the parsing effort and present the current state of the parser as the result to the caller), but this is not standard. With this encoding however, tail end recovery is implicit.

Funnily, some binary formats don't suffer from these limitations. Examples are raw audio, image and video encodings. So long as the encoder and decoder have previously agreed on resolution, any further bit in a video stream is simply the next bit used to create the current pixel, on the current scanline, in the current frame.

Granted, such formats have become rare in this day and age, and in audio and video we can under- or overflow the input buffer of the decoder in time.

And in images we can have an incomplete buffer. But what input we give it, cannot be 'illegal'. If we feed it random bits, the screen will show static, but it will be 'legal static'.

3 Design Principles

- Length can never be used in the way that it may be used to under- or overflow the message or message segment size.
- Our atom will be a bit, not a byte, as we cannot have any 'air' in our encoding.
- Complexity of the output must be 'like ASN.1, XML or JSON'. That is to say:
 - Encodes all eight-bit bytes and all lengths of those bytes.
 - Is complex in that it provides at least for strings, integers, floating points, 'hashtable' name-value lists, and arrays.
- We don't care about being truly efficient. The fact that we're binarily encoding should bring us enough of that. The main goal stands: any decoder input must produce an output.
- We don't care about deterministic reciprocity between encoding and decoding: if the input to the decoder produces a complex data structure, then the complex data structure, when fed to the encoder, doesn't need to produce a copy of the input.
- There exists a 'null' type (in fact, there are two). Which can be given explicitly, or we use it to fill in all sorts of blanks.
- On top, everything is implicitly inside a list. When the implicit list contains one element, it is a scalar. If you want to define a list with one element, you must explicitly define a list.
- At the end of input, however many zero bits as are required can be read. Since all reads are implicitly length delimited ('next byte' in string reading never pops more than nine bits, integers pop 64 bits, etc), this is a one-time affair only. All loops (strings, hashtables, arrays) take the end of input as to mean: jump out of the loop and return.

4 Specification

Explicitly written from the standpoint of the decoder. The encoder will use this specifiction simply in reverse.

- There are the following types (eight, so that they fit exactly into three bits):
 - NULL
 - Explicit NULL
 - BOOLEAN
 - INTEGER (64 bits signed)
 - FLOAT (of system double size)
 - STRING
 - ARRAY

– HASHTABLE

- The decoder starts assuming it has to fill a list. If, in the end, this list will contain exactly one element, it will return the element, not the list. In all other cases (empty list, list with more than one element), it will return the list.
- The decoder pops a type triplet and switches according to its value:
 - If it is a null, it will move to pop the next tuple.
 - If it is a boolean, it will pop the next bit to determine its value and move on to the next tuple.
 - If it is an int or a float, it will pop the next 64 bits, cast them to the machine representation, and move on to the next tuple.
 - If it is a string, it will go into a loop determining, per byte, whether or not to read it. It does this by popping a bit and if it is one, it will pop the byte. If it is zero, it assumes the string is finished and move on to the next tuple.
 - If it is a list, it does the same 'continuity bit popping' for each list element, and each element recurses into the list-tuple decoding.
 - If it is a hashtable, it pops a string (without type triplet popping) as key, and an element (as in the list), and does this in the same 'continuity bit popping' kind of way as with strings or list elements.
- Should an input end with one to seven zero bits, these should, per the specification above, be rendered as a list of one or more implicit NULL types (one null D-bit and three null type-bits, which can repeat a maximum of two times in seven bits), these will only not be discarded when they are explicit NULL types.

Note that implicit zeroes, enclosed in other relevant tuples, *are* absorbed as regular nulls.

5 Thoughts on Specific Forms of Sub-Encoding

5.1 Bigger type space encoding

Currently, the amount of types has been limited to eight. This happened to satisfy the requirements of encoding JSON. However, it may be that you're having a more extensive need for type encoding. In that case one could extend the type identifier sequence with one bit, giving you sixteen possible types. Add one more bit and you'll have 32, etc.

If then, however, you're not in need of all sixteen (or 32, or 64, or ...) types, you'll be forced to apply a modular arithmetic approach. Your type then becomes the modulo of type space available, divided by the amount of types needed. This will also make the encoder/decoder stages less deterministically linked (since you can have, in some cases, more than one discriminant for a type).

Alternatively, it may make sense to appreciate types as binary trees. This will take more processing, but it may also yield a more specific encoding, that won't necessarily require modular arithmetic. For example:

Produces the following JSON:

```
[ ]
```

Three implicit nulls are read from the input, which are all discarded. Since the implicit top-list now contains zero elements, it is presented as the result of the decoding.

The following input:

```
B: 0 0 1 0 0 0 0 0
```

Produces the following JSON:

```
null
```

An explicit null is read, followed by the read of two implicit nulls. The implicit nulls are discarded (because they are in the tail). Since the implicit top-list only contains one element, it is presented as a single null, not encapsulated in an array.

The following input:

```
B: 0 0 1 0 0 1 0 0
```

Produces the following JSON:

```
[ null, null ]
```

Here, two explicit (un-erasable) nulls are encoded in the implicit top-list, followed by two zero bits. The decoder reads these bits (and one more, which will also be virtually zero), and interprets that as an implicit null. The implicit null is then discarded.

7 Assessment

7.1 Advantages

- The format is ultimately streaming: the decoder never has to have more than `sizeof(double)` into its buffer, no 'looking back' is ever required.
- In spite of that, a complex structure a la JSON/XML/ASN.1 can be encoded.
- In fact, using the larger concept, *any* structure can be encoded.
- While the decoder *cannot fail*.
- Therefore, cannot fail mid-decoding either (you can receive partial messages and still have a valid encoding).

7.2 Disadvantages

- (In this implementation:) string encoding carries a penalty of 12.5% overhead in the limit.
- Encoding and decoding may not be deterministically linked: the same chaos may not be returned when passed through the decoder and encoder subsequently (the reverse order however, does deterministically link the encoder and the decoder).